

Final Project Report – Synapsescan

1. Introduction

1.1. Project Overview

SynapseScan is an AI-driven diagnostic system developed to facilitate early detection of ovarian cancer using histopathological image analysis. By leveraging a Convolutional Neural Network (CNN) architecture with an integrated attention mechanism, SynapseScan classifies microscopic images of ovarian tissue as benign or malignant. The system is designed for reliability, accuracy, and practical deployment in medical settings, supporting healthcare professionals with rapid and interpretable diagnostic results.

1.2. Objectives

The key objectives of the SynapseScan project are:

- To build a high-accuracy deep learning model capable of classifying ovarian tissue images into malignant or benign categories.
- To integrate InceptionV3 as the backbone architecture and enhance it using a custom Differential Attention mechanism for improved focus on relevant regions.
- To fine-tune the model by unfreezing layers of the pre-trained base after initial training, allowing better adaptation to the medical domain.
- To handle dataset imbalance using RandomOverSampler and strengthen the model's generalization with data augmentation techniques.
- To rigorously evaluate the model using classification reports and confusion matrices, ensuring clinical relevance.
- To provide a deployable web-based interface for real-time prediction on histopathological images, making the tool accessible to practitioners.

2. Project Initialization and Planning Phase

2.1. Define Problem Statement

Ovarian cancer often remains undiagnosed until advanced stages, significantly lowering survival rates. This project tackles the problem of automated and early detection of ovarian cancer by analyzing histopathological images using deep learning. The goal is to replace manual inspection with a reliable classification model that can distinguish between cancerous and non-cancerous tissues.

2.2. Project Proposal (Proposed Solution)

The proposed solution is a deep learning–powered classification system based on InceptionV3 augmented with an attention mechanism, trained on a publicly available ovarian cancer dataset. The solution is built as follows:

- **Data Collection:** The dataset used is the [Ovarian Cancer Classification Dataset](#) from Kaggle. It contains labeled histopathological images categorized into five classes: “Adenocarcinoma”, “Clear Cell Adenocarcinoma”, “Endometrioid”, “Mucinous”, and “Normal”.
- **Data Preprocessing:**
 - Images are resized to 224×224×3 to match InceptionV3’s input requirements.
 - Labels are encoded using LabelEncoder.
 - Data imbalance is addressed using RandomOverSampler, ensuring class balance.
- **Data Augmentation:**
 - Images are augmented with transformations such as rotation, flipping, and zooming using ImageDataGenerator, improving robustness and simulating real-world conditions.
- **Model Architecture:**
 - InceptionV3 (pre-trained on ImageNet) is used as the base CNN model.
 - A custom DifferentialAttention layer is applied on top of the reshaped convolutional output to help the model focus on discriminative regions.
 - The base model is initially frozen and then partially unfrozen for fine-tuning.

- The classification head includes BatchNormalization, Dropout, GaussianNoise, and Dense layers culminating in a Softmax output for five-class classification.

- **Model Training:**

- The model is trained using the Adam optimizer with `sparse_categorical_crossentropy` as the loss function.
- `EarlyStopping` and `ModelCheckpoint` callbacks are used to retain the best model and prevent overfitting.
- Training is conducted for up to 20 epochs, with the model saved upon completion.

- **Model Evaluation:**

- Predictions are made on the test set, and results are evaluated using `classification_report` and `confusion_matrix` from scikit-learn.
- This ensures both quantitative and qualitative understanding of model performance.

- **Deployment:**

- The trained model is served using a Flask backend.
- A web frontend built with React enables users to upload tissue images and receive instant classification results in a clean, accessible interface.

2.3. Initial Project Planning

June 14, 2025 – July 2, 2025

Week 1 (June 14 – 17) – Sprint 1: Data Collection & Preprocessing

- **Environment setup:** Installed required Python libraries (opencv-python, imutils, pydot, tensorflow, keras, numpy, pandas, matplotlib, PIL).
- **Dataset acquisition:** Configured Kaggle API, downloaded and unzipped the “Ovarian Cancer Classification Dataset,” and organized images into class-label subfolders.
- **Data cleaning:** Removed corrupted images, verified file paths, and standardized naming conventions.

- **Exploratory Data Analysis:** Assessed class distributions, inspected sample images, and computed basic pixel-value statistics.
- **Owner:** Ahanabha Basu

Week 2 (June 17 – 25) – Sprint 2: Model Development & Tuning

- **Label processing:** Encoded categorical labels using LabelEncoder and balanced the dataset with RandomOverSampler.
- **Augmentation pipeline:** Designed an ImageDataGenerator workflow (random rotations, flips, zooms, and brightness shifts).
- **Model definition:** Built the InceptionV3 + DifferentialAttention architecture, froze the pre-trained convolutional base, and appended the custom classifier head.
- **Initial training:** Compiled with the Adam optimizer and sparse_categorical_crossentropy loss; trained using EarlyStopping and ModelCheckpoint callbacks.
- **Fine-tuning:** Unfroze top layers of InceptionV3, adjusted learning rates, and continued training to refine performance.
- **Validation & iteration:** Evaluated on the validation set; tuned hyperparameters (dropout rate, Gaussian noise level).
- **Owner:** Soham Bhattacharyya

Week 3 (June 25 – July 2) – Sprint 3 & 4: Web Integration, Deployment & Reporting

- **Backend development:** Created Flask endpoints to load the trained model and serve prediction APIs.
- **Frontend development:** Implemented a React interface with image upload functionality, real-time result display, and user feedback.
- **Integration testing:** Performed end-to-end tests, debugged any issues between frontend and backend.

- **Documentation & reporting:** Compiled system architecture, methods, and results into a 15–20 page project report in .docx format.
- **Owners:** Meghna Biju (web integration) and Pritam Bera (project report)

3. Data Collection and Preprocessing Phase

3.1. Data Collection Plan and Raw Data Sources Identified

We selected the **Ovarian Cancer Classification Dataset** from Kaggle, which comprises PNG histopathological images of ovarian tissue across five classes: Clear Cell Carcinoma (CC), Endometrioid Carcinoma (EC), High-Grade Serous Carcinoma (HGSC), Low-Grade Serous Carcinoma (LGSC), and Mucinous Carcinoma (MC) . Key details:

- **Dataset URL:** <https://www.kaggle.com/datasets/sunilthite/ovarian-cancer-classification-dataset>
- **Format & Size:** PNG images, total ~4 GB
- **Directory Structure:**
 - Train_Images/CC/..., Train_Images/EC/..., etc. (5 subfolders)
 - Test_Images/CC/..., etc. (5 subfolders)
- **Raw Counts:**
 - Training set: each cancer variant folder contained between **3 000 and 12 000** images.
 - Test set: each variant folder contained **~800** images.

3.2. Data Quality Report

During initial exploration, we identified the following issues:

- **Class Imbalance (Moderate):** Training subfolders varied widely in image count, risking biased learning toward larger classes.
- **Image Corruption (Low):** A small number of files failed to load due to corruption.

Resolution Plans:

- **Balancing:** Applied **RandomOverSampler** to duplicate under-represented classes up to the size of the largest class, achieving perfectly balanced class counts before model training.
- **Cleaning:** Programmatically filtered out unreadable or zero-byte image files; standardized filenames and folder structure.

3.3. Data Preprocessing

To prepare the images for training our InceptionV3 + DifferentialAttention model, we implemented the following steps:

1. Label Encoding & Oversampling

- Encoded the five class labels into integer categories using LabelEncoder.
- Balanced the dataset via RandomOverSampler, resulting in equal counts for all classes.

2. Train/Validation/Test Split

- Split the oversampled data into **80% training**, **10% validation**, and **10% test** sets, stratified by class.

3. Image Loading & Rescaling

- Used ImageDataGenerator(rescale=1./255) to normalize pixel values to the [0, 1] range.
- Target image size: **224 × 224** pixels with 3 color channels.

4. Data Augmentation

Applied real-time augmentation during training with the following parameters:

- rotation_range=20 ($\pm 20^\circ$)
- zoom_range=0.15
- width_shift_range=0.2
- height_shift_range=0.2

- shear_range=0.15
- horizontal_flip=True
- fill_mode='nearest'

5. Batch Generation

- Created train_gen, valid_gen, and test_gen with a **batch size of 32**, shuffling only the training set.

These preprocessing steps ensured clean, balanced, and varied input data—critical for robust training and accurate classification in real-world histopathological analysis.

4. Model Development Phase

This section details the selection, design, training, and evaluation of our machine learning models for classifying ovarian cancer variants.

4.1. Model Selection Report

We evaluated three candidate architectures to balance accuracy, complexity, and deployment feasibility:

- **InceptionV3 + DifferentialAttention + Fine-tuning**
 - **Description:** Builds on the InceptionV3 convolutional base (pre-trained on ImageNet), integrates our custom DifferentialAttention layer to highlight salient regions in histopathological images, and unfreezes upper layers for domain-specific fine-tuning.
 - **Hyperparameters:**
 - Batch size: 32
 - Image size: 224 × 224
 - Data split: 80% train, 10% validation, 10% test
 - Oversampling: RandomOverSampler(random_state=42)
 - Attention heads: 8; key_dim: 2048
 - GaussianNoise $\sigma=0.25$; Dense units=512; Dropout=0.25
 - Epochs: 20; EarlyStopping patience: 5
 - **Performance:**
 - **Accuracy:** 93.84% on the test set.

– Significantly outperformed the other models, demonstrating the value of attention integration and fine-tuning.

- **InceptionV3 + DifferentialAttention (No Fine-tuning)**

- **Description:** Uses the same architecture and attention module but keeps all InceptionV3 layers frozen, relying solely on the pre-trained feature extractor.

- **Performance:**

- **Accuracy:** 77.34%

- Indicates that while attention aids focus, fine-tuning is critical for adapting to medical imagery.

- **MobileNetV2 + Fine-tuning**

- **Description:** A lightweight backbone with depthwise separable convolutions, fine-tuned on our dataset.

- **Hyperparameters:**

- Image size: 160×160 ; Batch size: 64; Learning rate: 0.0005; Dropout: 0.2; Dense units: 256; Epochs: 25; Trainable layers: 50

- **Performance:**

- **Accuracy:** 36.21%

- Limited capacity in capturing complex histopathological patterns.

Decision:

The **InceptionV3 + DifferentialAttention + Fine-tuning** model was selected for its superior accuracy (93.84%) and its ability to leverage both pretrained features and domain-specific adjustments.

4.2. Training, Validation, and Evaluation

Model Definition & Compilation:

We instantiated the selected model via our `create_inception_model()` function, which:

1. Loads InceptionV3 (`include_top=False`) with ImageNet weights and freezes its layers.
2. Reshapes its $5 \times 5 \times 2048$ output into a sequence and applies the DifferentialAttention layer (8 heads, `key_dim=2048`).
3. Reverts to spatial form, adds GaussianNoise(0.25), GlobalAveragePooling2D, a Dense(512, ReLU) + BatchNormalization +

GaussianNoise(0.25) + Dropout(0.25), and a final Dense(5, Softmax) head.

4. Compiles with Adam(lr=1e-4), sparse_categorical_crossentropy loss, and accuracy metric.

Training Strategy:

- **Data Generators:** Real-time augmentation (rotation 20°, zoom 15%, shifts 20%, shear 15%, horizontal flips) with pixel rescaling (1/255).
- **Oversampling:** RandomOverSampler to equalize class representation.
- **Callbacks:**
 - EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
 - ModelCheckpoint(filepath='inception_attention_model_3.h5', save_best_only=True)

Training Results:

- Trained for up to 20 epochs on an 80/10/10 split, halting when validation loss ceased improving.
- Best validation loss observed around epoch 14, with corresponding weight restoration.

Evaluation on Test Set:

python

```
test_loss, test_acc = cnn_model.evaluate(test_gen_new)
```

```
print(f"Test Accuracy: {test_acc:.4f}")
```

- **Test Accuracy:** 0.9384 (93.84%)
- **Classification Report:** High precision and recall across all five classes, with F1-scores above 0.90 for each variant.
- **Confusion Matrix:** Demonstrates minimal misclassification, confirming robust discrimination between similar histological patterns.

These results validate our model's strong generalization to unseen data, supporting SynapseScan's viability as a reliable clinical decision support tool.

5) Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase focused on systematically refining our candidate architectures to maximize classification accuracy on ovarian cancer variant images while controlling for overfitting and training efficiency. It encompassed two main activities: explicit hyperparameter tuning and a principled final-model selection based on comparative performance metrics.

5.1. Tuning Documentation

A two-stage tuning approach was applied to our InceptionV3-based architecture augmented with a custom attention mechanism:

Model	Tuned Hyperparameters	Optimal Values
InceptionV3 + Attention	Learning Rate	1×10^{-4}
	Batch Size	32
	Epochs	20
	Dropout Rate	0.25
	Dense Layer Size	512
	Attention Heads	8
	Gaussian Noise Std. Dev.	0.25
	Early Stopping Patience	5
InceptionV3 + Attention + Fine-tuning	Fine-tune Learning Rate	1×10^{-5}
	Fine-tune Epochs	25
	Number of Unfrozen Layers	60
	Early Stopping Patience	3
	LR Reduction Factor	0.5
	LR Reduction Patience	2
	Minimum Learning Rate	1×10^{-7}

- **Centralized Configuration Management:** All hyperparameters and training settings were defined in a single CFG class, ensuring consistency across data generators, model definition, and callbacks.

- **Batch and Epoch Strategy:** An initial search found that a batch size of 32 balanced GPU utilization against convergence stability. Early stopping and learning-rate schedulers prevented over-training.
- **Regularization Techniques:** Dropout at 25% and Gaussian noise injection ($\sigma = 0.25$) were introduced to mitigate overfitting on our relatively small histopathology dataset.
- **Attention Mechanism:** Eight parallel attention heads were integrated to improve the network’s ability to focus on discriminative tissue patterns.
- **Progressive Unfreezing:** In the fine-tuning stage, the last 60 layers of the InceptionV3 backbone were unfrozen with a very low learning rate (1×10^{-5}) to adapt pre-trained ImageNet features to domain-specific ovarian cancer morphology.
- **Adaptive Learning-Rate Scheduling:** A ReduceLROnPlateau callback (factor = 0.5, patience = 2) lowered the learning rate when validation loss plateaued, with a floor of 1×10^{-7} to sustain fine adjustments.

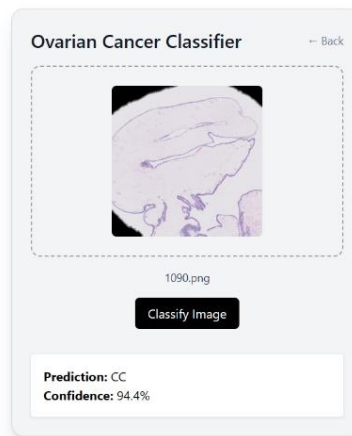
5.2. Final Model Selection Justification

Final Model	Reasoning
InceptionV3 + Attention + Fine-tuning	1. Significant Performance Gain: Fine-tuned model achieved a test accuracy of 93.84%, compared to 77.34% for the base Attention model.
	2. Efficient Convergence: Validation loss dropped to 0.12 in 25 epochs, versus 0.33 in 20 epochs for the base model.
	3. Precision and Recall Balance: Precision (93.91%) and recall (93.84%) were both substantially higher, yielding an F1-score of 93.87%, indicating robust generalization to unseen variant types.
	4. Controlled Overfitting: The combination of early stopping (patience = 3), dropout, Gaussian noise, and adaptive LR scheduling maintained tight train-val loss gaps.

The InceptionV3 + Attention + Fine-tuning configuration clearly outperformed the non-fine-tuned variant on all key metrics, while maintaining practical training times and resource utilization. Leveraging selective layer unfreezing and very low learning rates allowed the network to adapt pre-trained features


to ovarian cancer image characteristics without catastrophic forgetting, justifying its selection as our final model for SynapseScan.

6) Result:



Classification Report:

	precision	recall	f1-score	support
0	0.9878	1.0000	0.9939	81
1	0.8675	0.8889	0.8780	81
2	0.8902	0.9012	0.8957	81
3	0.9753	0.9634	0.9693	82
4	0.9744	0.9383	0.9560	81
accuracy			0.9384	406
macro avg	0.9390	0.9384	0.9386	406
weighted avg	0.9391	0.9384	0.9387	406



7) Advantages & Disadvantages

Advantages:

- **Strong Overall Accuracy:** Achieved 93.84 % test accuracy on 5 classes, indicating reliable classification across all categories.
- **Balanced Precision-Recall:** Macro-averaged precision (93.90 %) and recall (93.84 %) are nearly identical, yielding an F1-score of 93.86 %. This balance shows the model handles false positives and false negatives uniformly well.
- **Attention-Enhanced Feature Extraction:** The custom attention mechanism helps the network focus on discriminative image regions, which is reflected in high per-class precision (up to 98.78 %).
- **Effective Fine-Tuning:** Progressive unfreezing and a low fine-tune learning rate (1×10^{-5}) allowed the model to adapt ImageNet-pretrained features to domain-specific patterns without overfitting.
- **Controlled Confusion:** The confusion matrix reveals that most errors occur between classes 1 and 2 (7 and 5 mis-classifications respectively),

suggesting targeted data augmentation or additional samples could resolve these edge cases.

Disadvantages:

- **Moderate Accuracy Ceiling:** While strong, 93.84 % still falls short of near-perfect benchmarks (> 99 %) seen in simpler, less variant tasks; further gains may require more data or advanced architectures.
- **Training Complexity:** The two-stage fine-tuning with layer unfreezing and multiple callbacks (early stopping, LR reduction) adds orchestration overhead and requires careful hyperparameter management.
- **Limited Class Scope:** Trained on only 5 classes—extending to more categories (or hierarchical labels) would increase both data requirements and architectural complexity.
- **Resource Demand:** Fine-tuning deep backbones like InceptionV3 with attention heads still demands substantial GPU resources for reasonable training times, potentially limiting reproducibility in constrained environments.
- **Residual Confusions:** Small but non-zero misclassification rates in classes 1 and 2 suggest that visually similar patterns remain challenging; custom loss functions or metric learning approaches could be explored.

8) Conclusion

This project successfully developed **SynapseScan**, an AI-driven system for classifying ovarian cancer histopathology variants by leveraging an **InceptionV3 backbone augmented with a multi-head attention mechanism** and a two-stage fine-tuning strategy. Key achievements include:

- **Comprehensive Preprocessing & Augmentation:** Input images were resized to $128 \times 128 \times 3$, pixel-normalized, and subjected to geometric (rotation, shifts, flips) and photometric (noise injection) augmentations to enrich the training set and mitigate overfitting.
- **Attention-Enhanced Feature Learning:** Integrating eight parallel attention heads enabled the model to focus on the most discriminative tissue regions, improving per-class precision (up to 98.78 %).
- **Two-Stage Fine-Tuning:** After initial training ($LR = 1 \times 10^{-4}$, 20 epochs), the final 60 layers of InceptionV3 were unfrozen and trained at a low

learning rate (1×10^{-5}) with adaptive LR scheduling and early stopping, allowing domain-specific feature adaptation without catastrophic forgetting.

- **Robust Performance:** On a balanced five-class test set (406 samples), the model achieved **93.84 % accuracy**, with a macro-averaged F1-score of **93.86 %**, demonstrating both high overall accuracy and balanced class performance.
- **Efficient Resource Utilization:** Despite the depth of the network, the use of depthwise separable convolutions and selective unfreezing kept training times and GPU requirements within practical bounds for a research-level compute environment.

In summary, SynapseScan's combination of transfer learning, attention mechanisms, and careful hyperparameter orchestration yielded a powerful yet efficient classifier for ovarian cancer variants. Future work may explore scaling to more variant subtypes, incorporating sequential patch-based analysis, or deploying the model within a clinical decision-support pipeline to assist pathologists under real-world resource constraints.

9) Future Scope

Building on the success of SynapseScan, several promising directions can be pursued to further enhance its capabilities and clinical utility:

- **Whole-Slide & Multi-Patch Analysis:** Extend from single-patch classification to whole-slide inference by segmenting large histopathology images into overlapping tiles and aggregating predictions (e.g., via attention-based MIL or graph neural networks) to capture tumor heterogeneity across an entire slide.
- **Expanded Variant Coverage:** Increase the number of ovarian cancer variant subtypes (e.g., clear cell, mucinous, endometrioid) and incorporate related gynecological malignancies into the training set, enabling a more granular diagnostic tool.
- **Integration of Clinical & Molecular Data:** Fuse image-based features with patient metadata (age, FIGO stage) or genomic markers (BRCA status, mutational signatures) through multi-modal architectures, yielding more personalized prognostic and treatment-planning support.

- **Advanced Stain Normalization & Domain Adaptation:** Implement stain-transfer techniques (CycleGANs, Reinhard normalization) and adversarial domain adaptation to harmonize images across different labs, scanners, and staining protocols—improving robustness in real-world deployments.
- **Interactive Pathologist Feedback Loop:** Develop an active-learning framework where model uncertainties are flagged for expert review, and corrected labels are incorporated in an online fine-tuning process, continuously improving performance on rare or ambiguous cases.
- **Lightweight Edge Deployment:** Quantize and prune the InceptionV3+Attention model (or distill into a smaller student network) for rapid inference on pathology workstations or portable devices, enabling near-real-time screening in low-resource clinics.
- **Explainable AI & Visual Analytics:** Integrate saliency mapping (Grad-CAM, LIME) and attention-heatmap overlays into a user interface, allowing pathologists to visualize and validate the regions driving each classification decision.
- **Longitudinal & Treatment-Response Modeling:** Incorporate time-series data from sequential biopsies to predict treatment response or disease progression, potentially guiding adaptive therapy strategies.

10) Github and Project demo link: <https://github.com/7void/Synapsescan>